

Side Channels, and Other Gaps Between Theory and Practice

PAUL KOCHER

Korea Crypto Forum
Nov 15, 2019 9:00-10:30



*If the surgery proves unnecessary, we'll
revert your architectural state at no charge.*

Intro

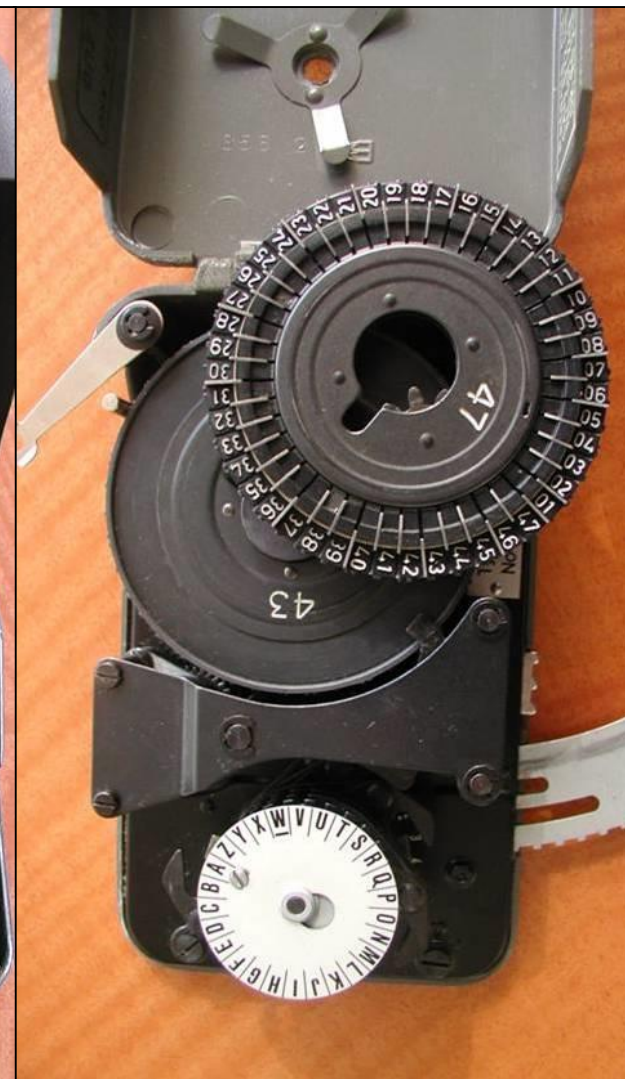
Researcher & entrepreneur

Technical work

- Protocols (incl. SSL v3.0 / TLS “🔒”)
- Timing attacks/side channels
- Differential power analysis & countermeasures
- Numerous HW/ASIC projects
 - Pay TV, anti-counterfeiting, keysearch
- Renewability & forensics
 - Blu-ray BD+, Vidity/SCSA...
- CryptoManager ASIC & mfg solutions
- Spectre Attacks
 - Co-discovery w/Jann Horn of Google

Non-technical

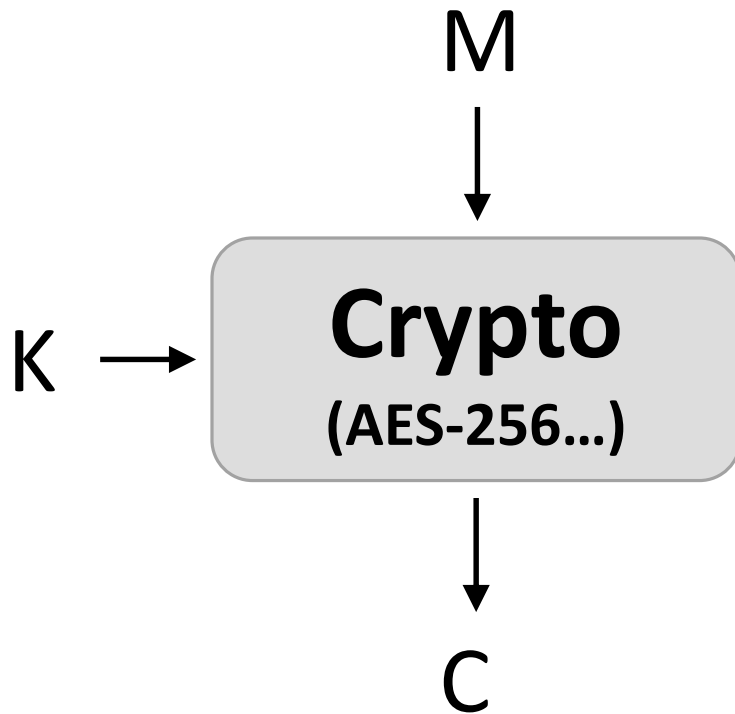
- Founded and led Cryptography Research (1995-April 2017)
 - Bootstrapped – no external funding
 - Multiple stage evolution:
Consulting → Licensing → Products → → Solutions
 - Acquired by Rambus (\$342.6M)
- Co-founded ValiCert (IPO 2001, acquired 2003)
- Advisor+investor to security start-ups
- Favorite nonprofits include Doctors Without Borders (MSF), Muttville, IACR, Bay Area Discovery Museum



Historical cipher machines:

Algorithm limited by mechanical constraints
Security limited by algorithm complexity

Modern age: Algorithm designers have won



More compute favors the cryptographer over the cryptanalyst

- 2-8X CPU power = 2X key length = **square** the effort for cryptanalysis
- Example: DES to triple DES = 3X CPU, keysearch from $\sim 2^{56}$ \rightarrow $\sim 2^{112}$

Modern crypto algorithms generally have LARGE safety margins:

Survive very conservative scenarios

e.g., exabytes of adaptively-chosen plaintexts...

Survive huge adversary compute power

e.g., attacker with more than world's compute power

Global bitcoin mining $\sim 2^{91}$ hashes/year vs AES brute force $\sim 2^{255}$

$2^{255}/2^{91} = 23,384,026,197,294,446,691,258,957,323,460,528,314,494,920,687,616$ years

Defenders have defined a “game” in which attackers have no winning strategy!

Attacker options

Attacker won't play the according to the defender's rules (i.e. “cheat”)

- e.g., attack gap between defender's assumptions & reality

Invest in cryptanalysis research to find easier attack:

- Success unlikely (= breakthrough research)
- “Success” = probably not practical (computational complexity, input needs...)

Invest in faster computers, e.g. to do brute force:

- Wildly impractical using known technology
- Quantum computer research?
 - Expensive/uncertain research, many years until anything practical, and defenders just upgrade

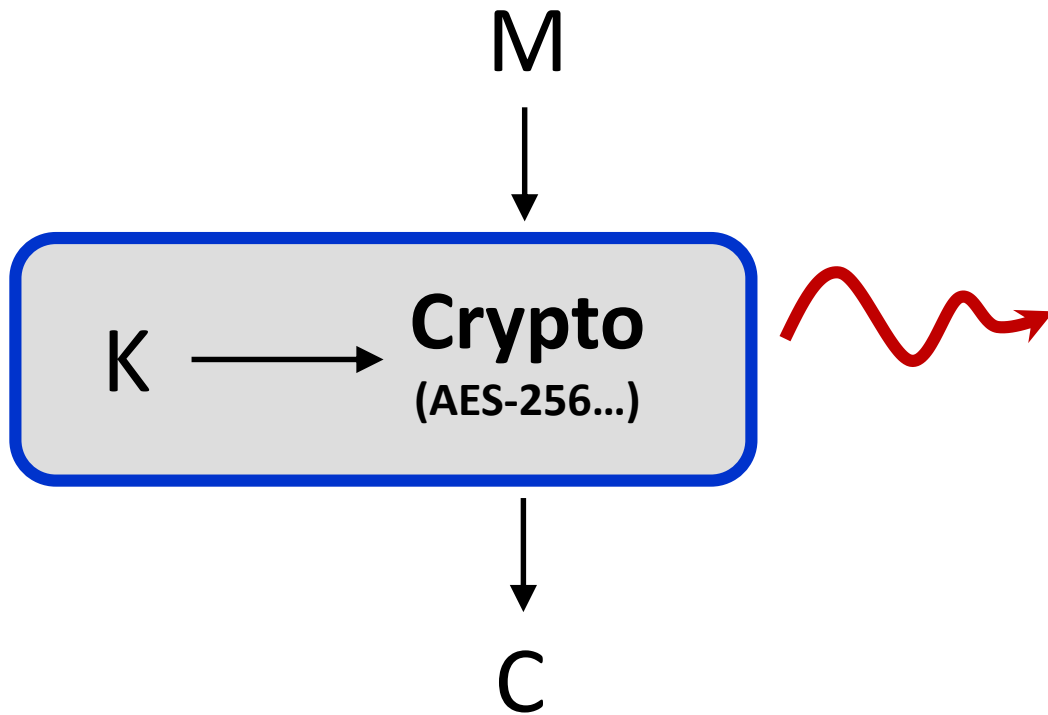
Give up

- Defender wins

Model (algorithm) vs. reality (implementation)

Algorithm assumptions:

Attacker sees many plaintexts
and/or ciphertexts, but nothing else



Practice:

Implementation uses semiconductors
(and maybe software), not pure math

- Timing variations
 - Power consumption
 - RF emissions
 - Heat
 - Debug registers
 - Microscope images of a chip
 - Erroneous outputs
 - Error messages
- etc

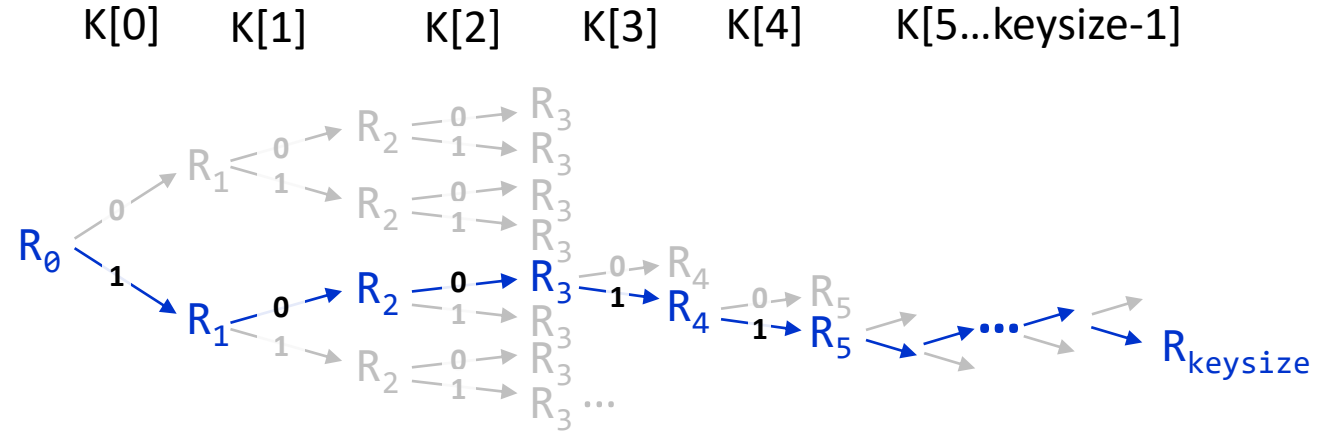
Exploiting tiny gaps between assumptions & reality

Timing

```

Let  $R_0$  = input message  $M$ 
for  $i$  in  $0..keysize-1$ :
    Let  $R_{i+1} = F(R_i, K[i])$ 
Output  $R_{keysize}$ 

```



Suppose:

- ▶ The time t_i to compute $F(\cdot)$ in iteration i varies slightly depending on its inputs (R_i and bit $K[i]$)
- ▶ Adversary can measure the total operation time (all iterations + some measurement error)

Each measured time T reflects:

$$T = \text{constant} + \text{noise} + \sum_{i=0}^{keysize-1} t_i \quad \leftarrow \text{time to compute each } F(R_i, K[i])$$

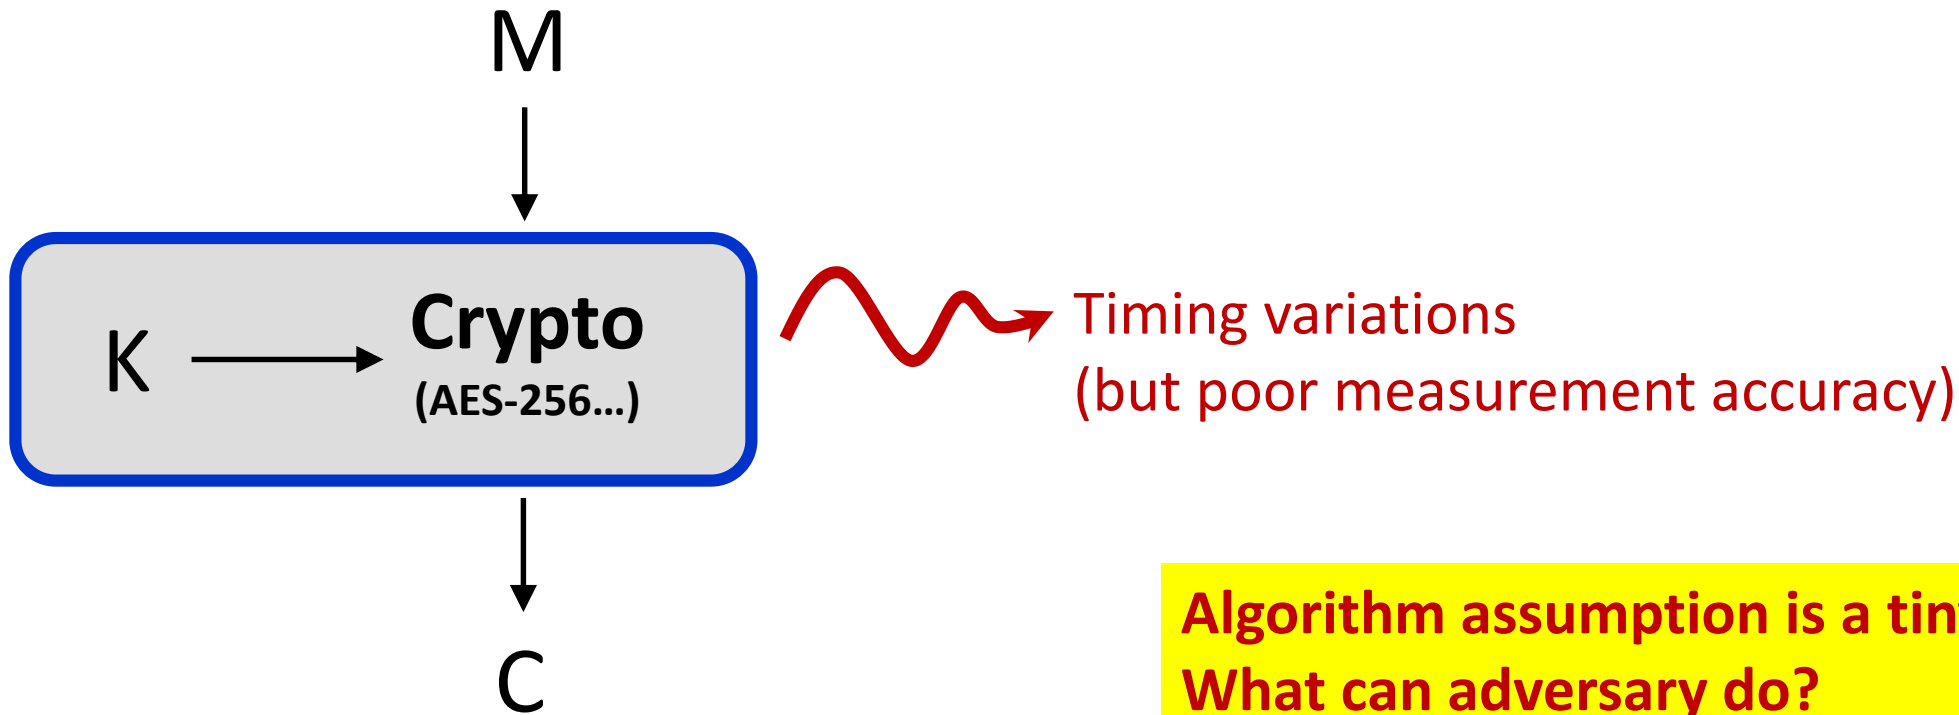
Loop overhead...

Measurement error, timing jitter...

Does it matter?

Algorithm assumptions:

Attacker sees many plaintexts
and/or ciphertexts, **but nothing else**



**Algorithm assumption is a tiny bit wrong.
What can adversary do?**

Attack inputs: n messages ($M_{1..n}$) and their timings ($T_{1..n}$)

Attack goal: Given $K[0..i-1]$, find the next bit $K[i]$

Attacker knows $K[0..i-1]$, so can find R_i for each M

Attacker doesn't know $K[i]$, so doesn't know if:

$$\begin{aligned} R_{i+1} &= \mathbf{F}(R_i, 0) \quad \text{or} \\ R_{i+1} &= \mathbf{F}(R_i, 1) \end{aligned}$$

```
Let  $R_0$  = input message  $M$ 
for  $i$  in  $0..keysize-1$ :
    Let  $R_{i+1} = \mathbf{F}(R_i, K[i])$ 
Output  $R_{keysize}$ 
```

For each M , attacker can estimate how long victim would take to compute $\mathbf{F}()$ for each key bit value:

$$\begin{aligned} t_{i,0} &= \text{time to compute } \mathbf{F}(R_i, 0) && \leftarrow \text{if } K[i]=0 \\ t_{i,1} &= \text{time to compute } \mathbf{F}(R_i, 1) && \leftarrow \text{if } K[i]=1 \end{aligned} \quad \left. \vphantom{\begin{aligned} t_{i,0} \\ t_{i,1} \end{aligned}} \right\} \text{One of these this is a (tiny) part of the overall time } T$$

The observed time is:

$$T = \text{constant} + \text{noise} + t_0 + t_1 + t_2 + \dots + t_i + \dots + t_{keysize-1}$$

\swarrow Observed time for a message

\nwarrow Either $t_{i,0}$ or $t_{i,1}$

Correlation

$$T = \text{constant} + \text{noise} + t_0 + t_1 + t_2 + \dots + t_i + \dots + t_{\text{keysize}-1}$$

Observed time for a message

Either $t_{i,0}$ or $t_{i,1}$

If $\mathbf{F}(R_i, 0)$ was actually computed, T will be correlated to $t_{i,0}$ over messages $M_{1..n}$

If $\mathbf{F}(R_i, 1)$ was actually computed, T will be correlated to $t_{i,1}$ over messages $M_{1..n}$

Simple attack to find key bit i :

- Predict $t_{i,0}$ and $t_{i,1}$ for each $M_{1..n}$
- For M where $t_{i,0}$ is faster than average, how many also have T faster than average? Votes for $K[i] = 0$
- For M where $t_{i,1}$ is faster than average, how many also have T faster than average? Votes for $K[i] = 1$

Given enough messages, the actual value for $K[i]$ will get more votes

Overcoming noise

How many timing measurements do we need?

- σ_1 = noise = standard deviation among T (measurement error, unknown t_i)
- σ_2 = signal = average difference between $t_{i,0}$ and $t_{i,1}$
- Need $N \propto \left(\frac{\sigma_1}{\sigma_2}\right)^2$ measurements

Example: $\sigma_1 = 1\mu\text{s}$ variation, $\sigma_2 = 1\text{ms}$, $N \approx 10^6$

Noise 1000X signal

- One set of measurements only (reuse for all key bits)
- Error detection: if a bit is guessed wrong, subsequent correlations go away (= back up + fix)

More noise = attack still works (but need more measurements)

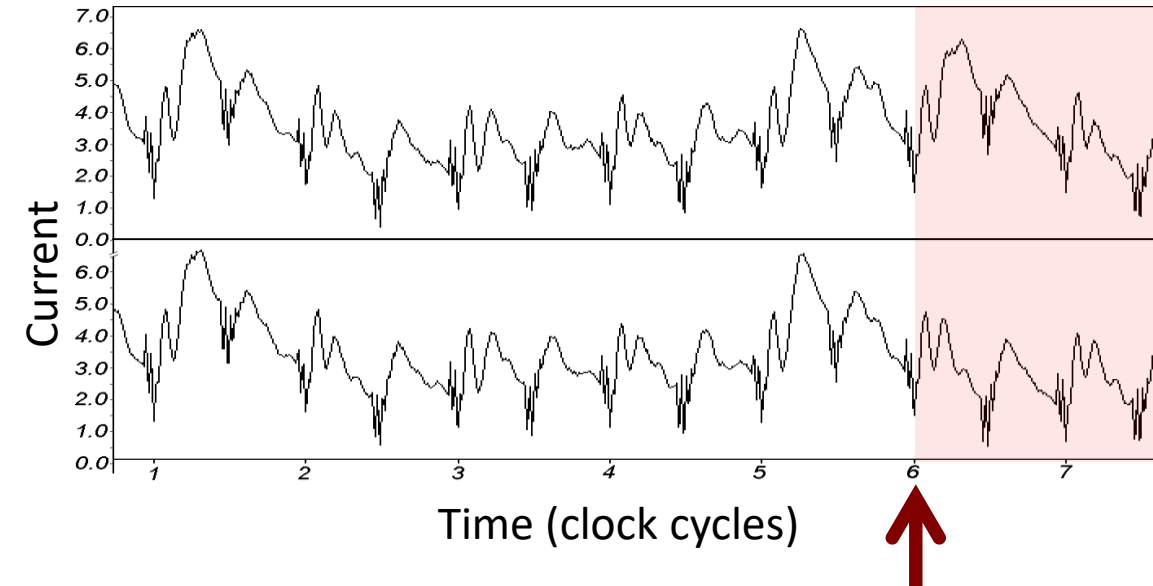
Generally not sufficient to add modest random delays or limits on timer resolution

Simple Power Analysis

Attackers can measure more than transaction time

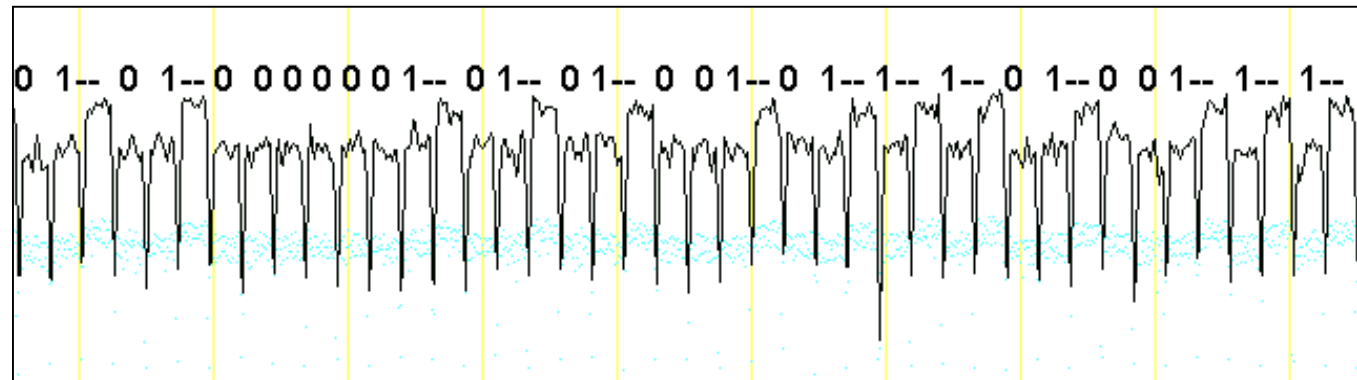
Example: Power consumption or RF

- Easy to measure if physically near device
- High bandwidth (many measurements during computation)



Simple case: power consumption is visibly different for $K[i]=0$ vs $K[i]=1$

```
for i in 0..keysize-1:  
    Let  $R_{i+1} = F(R_i, K[i])$ 
```



Signal-to-noise ratios for power & RF can also be poor

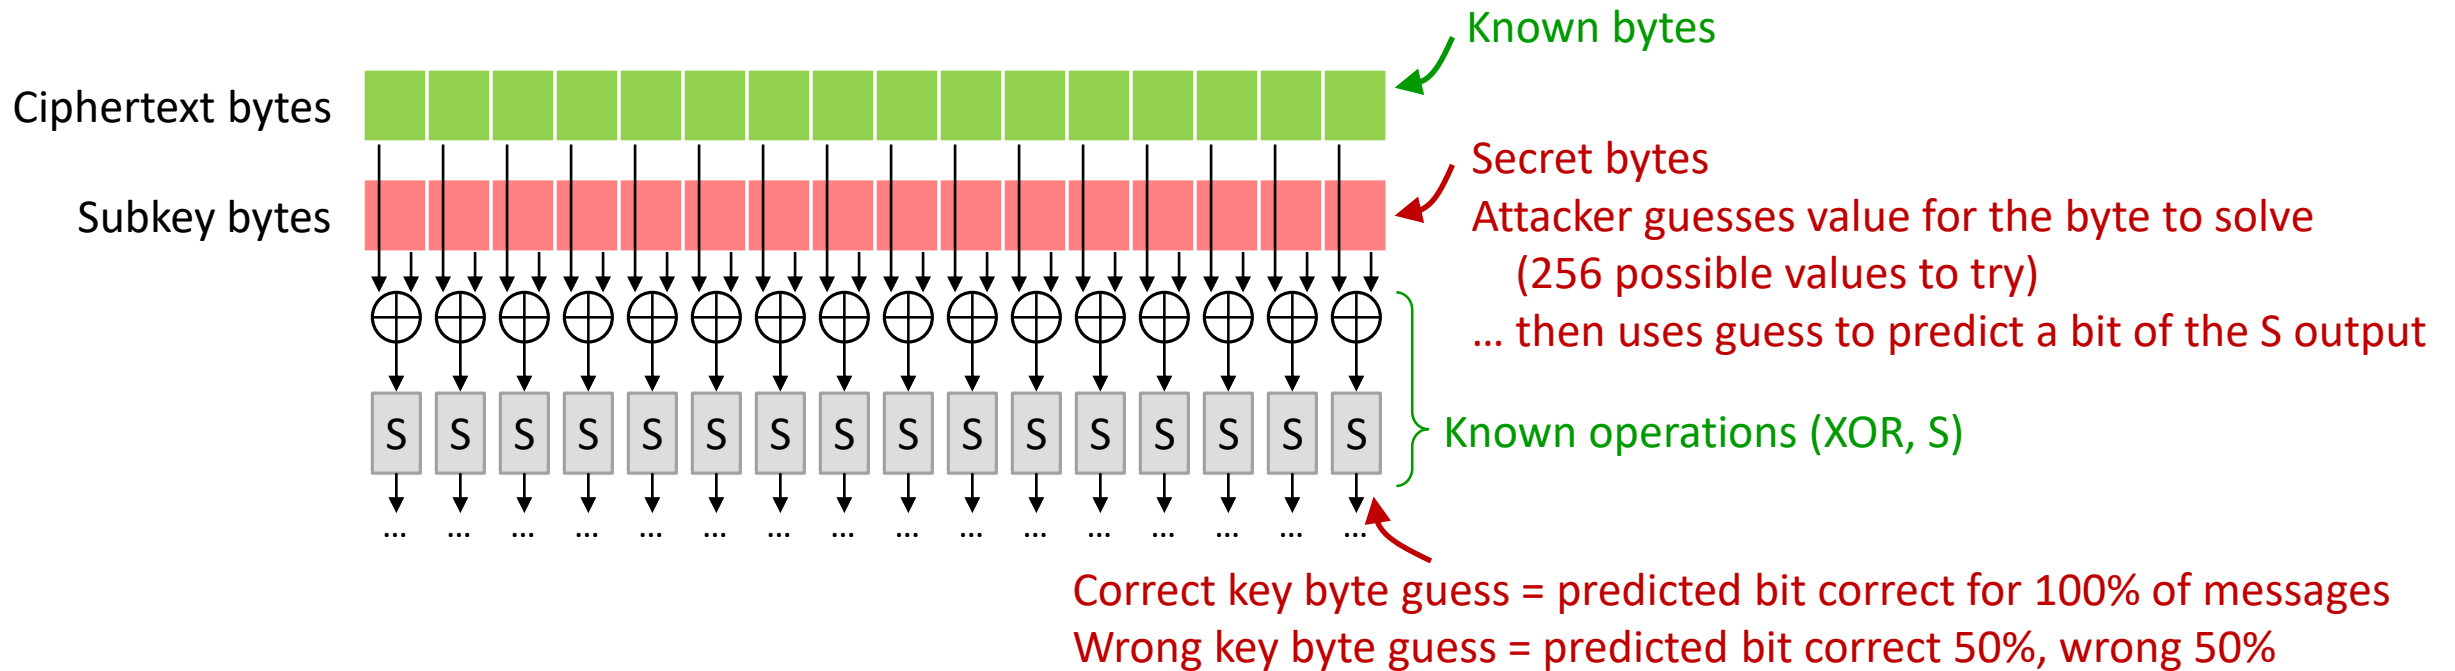
- Signals of interest may just be a few logic gates/wires in a large chip

Can these noisy power measurements be leveraged using statistics?

Overcoming noise: AES decryption

Attack scenario: AES decryption, where adversary can observe ciphertexts & power measurements

In first decryption round of AES:



Overcoming noise

Input message (ciphertext)

399632A62CA005E77710E3B947B1B65F

C06854D3EA68B82B8BFAADF73154ED17

01CC39A29605DD66A881D5128AC7FFA5

8FFC2EF2E831EF098514CA537F95EFA0

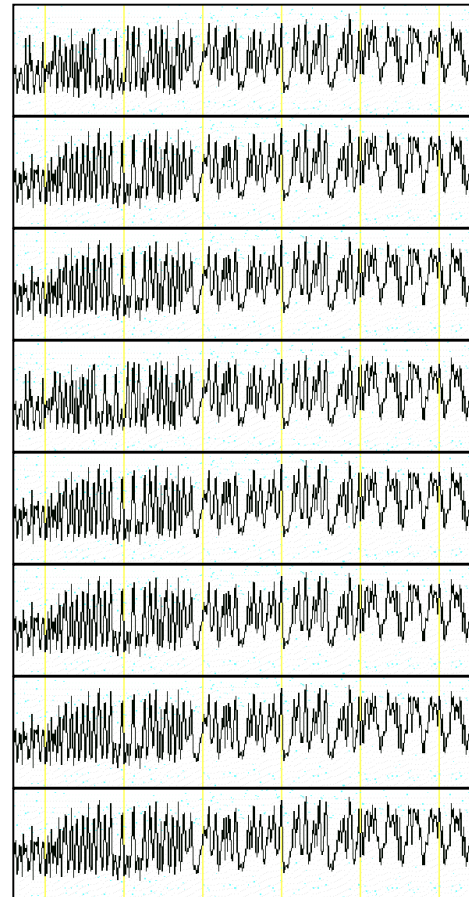
F02C294367605EB6ED393275C280587F

EEEE63394CE9DF029C52CCC9D69D8834

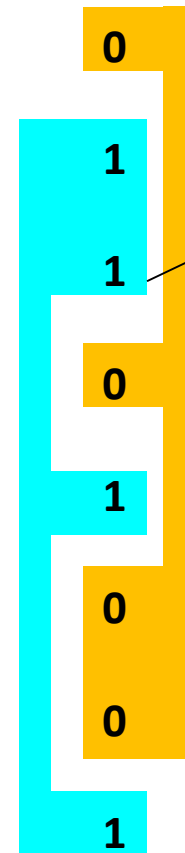
621E0462C66D4609AC0318A0284E1B7F

2FCA8B635D3DABBD8FEFC2AD7694D456

Power trace

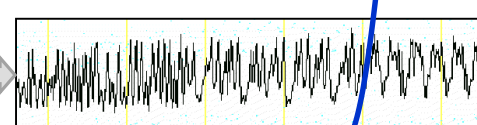
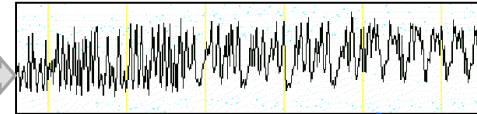


Predicted S output bit

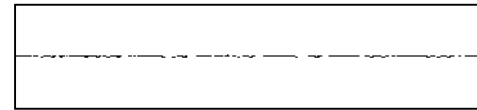


Compute average of
traces predicting 0

Compute average of
traces predicting 1

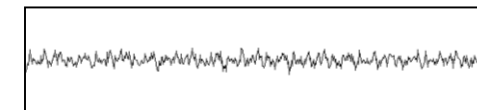


subtract

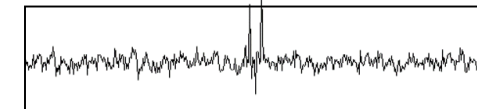


Rescale vertical

Incorrect key byte guess



Correct key byte guess



DPA: Overcoming noise

10000 traces:

Input message (ciphertext)

Power trace

Predicted
S output bit

Incorrect guess:

- Predictions are essentially random
- No significant differences in averages (just sampling noise)
- More messages/traces = subtraction result converges to 0 (flat)

Correct guess for key byte:

- Assignment of power traces reflects actual gate outputs (on/off)
- Difference in averages = effect in AES circuitry + sampling noise
- More inputs = less noise, clearer signal

Compute average of
traces predicting 0

Compute average of
traces predicting 1

subtract

Rescale vertical

Incorrect key byte guess

Correct key byte guess

Sharing hardware

Power & timing attacks assumed an external adversary

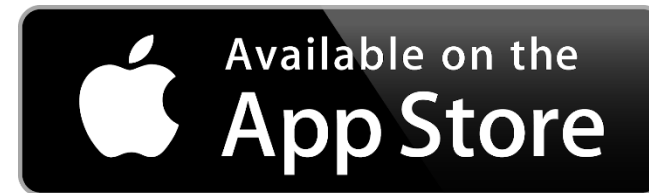
But what if attackers can run on the victim's computer...?

- Cloud
- Mobile
- Browser



JavaScript

WebAssembly



Transmitters & receivers

Shared resources => channels

- Program A uses resource -> program B observes use of resource

Transmitter: Victim code leaks sensitive data into channel

- Natural result of using system resources

Receiver: Adversary code reads from the channel

- Analyzes measurements to infer victim secret

Examples of side channel transmitters & receivers

| Channel | Transmit: Secret affects victim's... | Receive: Attacker measures... |
|------------------|--------------------------------------|---|
| CPU utilization | CPU usage | how much time attacker gets |
| DRAM usage | DRAM allocations | free memory |
| DRAM bandwidth | DRAM access rate | memory performance |
| cache state | memory accesses pattern | cache additions/evictions (see next slides) |
| GPU time | use of GPU | GPU availability |
| thermal | heat generation | effects of temperature on CPU speed |
| swap space | size of (compressed?) swap | disk/swap usage or free space |
| RNG seed pool | timing/amount of random data us | RNG latency (e.g. rand instruction) |
| disk utilization | timing/amount of file read/writes | disk/file system performance |
| signal coupling | crosstalk to analog circuitry | LSBs from A/D converters, oscillators... |

... many others: acoustic effects, LED brightness, debug/performance register state...

Microprocessor performance

CPU market driven has been by speed

$$\text{Time} = (\text{clock rate}) \times (\# \text{ clock cycles})$$


Rates grew fast to ~4 GHz... then plateaued:

- 1990: 80486 at 33MHz
 - 2004: Pentium 4 at 3.8 GHz
 - 2018: i7-8086K 'up to' 5.0 GHz
- } +40.36% / year
- } +1.98% / year

Performance gains => reducing # clocks

- Getting more done per clock

Moore's Law allows huge complexity

- Squeeze gains from cumulative effect of many complex optimizations

Getting more done per clock

Program = series of instructions

- Arithmetic/logic operations (add, multiply, divide...)
- Memory operations (load, store)
- Flow control (jump = decide which instruction to do next)

Example operation “ $R1 = R2 + R3$ ”

Fetch instruction from memory
Decode instruction
Find register inputs R2,R3
Perform add
Store result

Simple CPU

~200 clock cycles

~1 clock cycle

~1 clock cycle

~1 clock cycle

~1 clock cycle

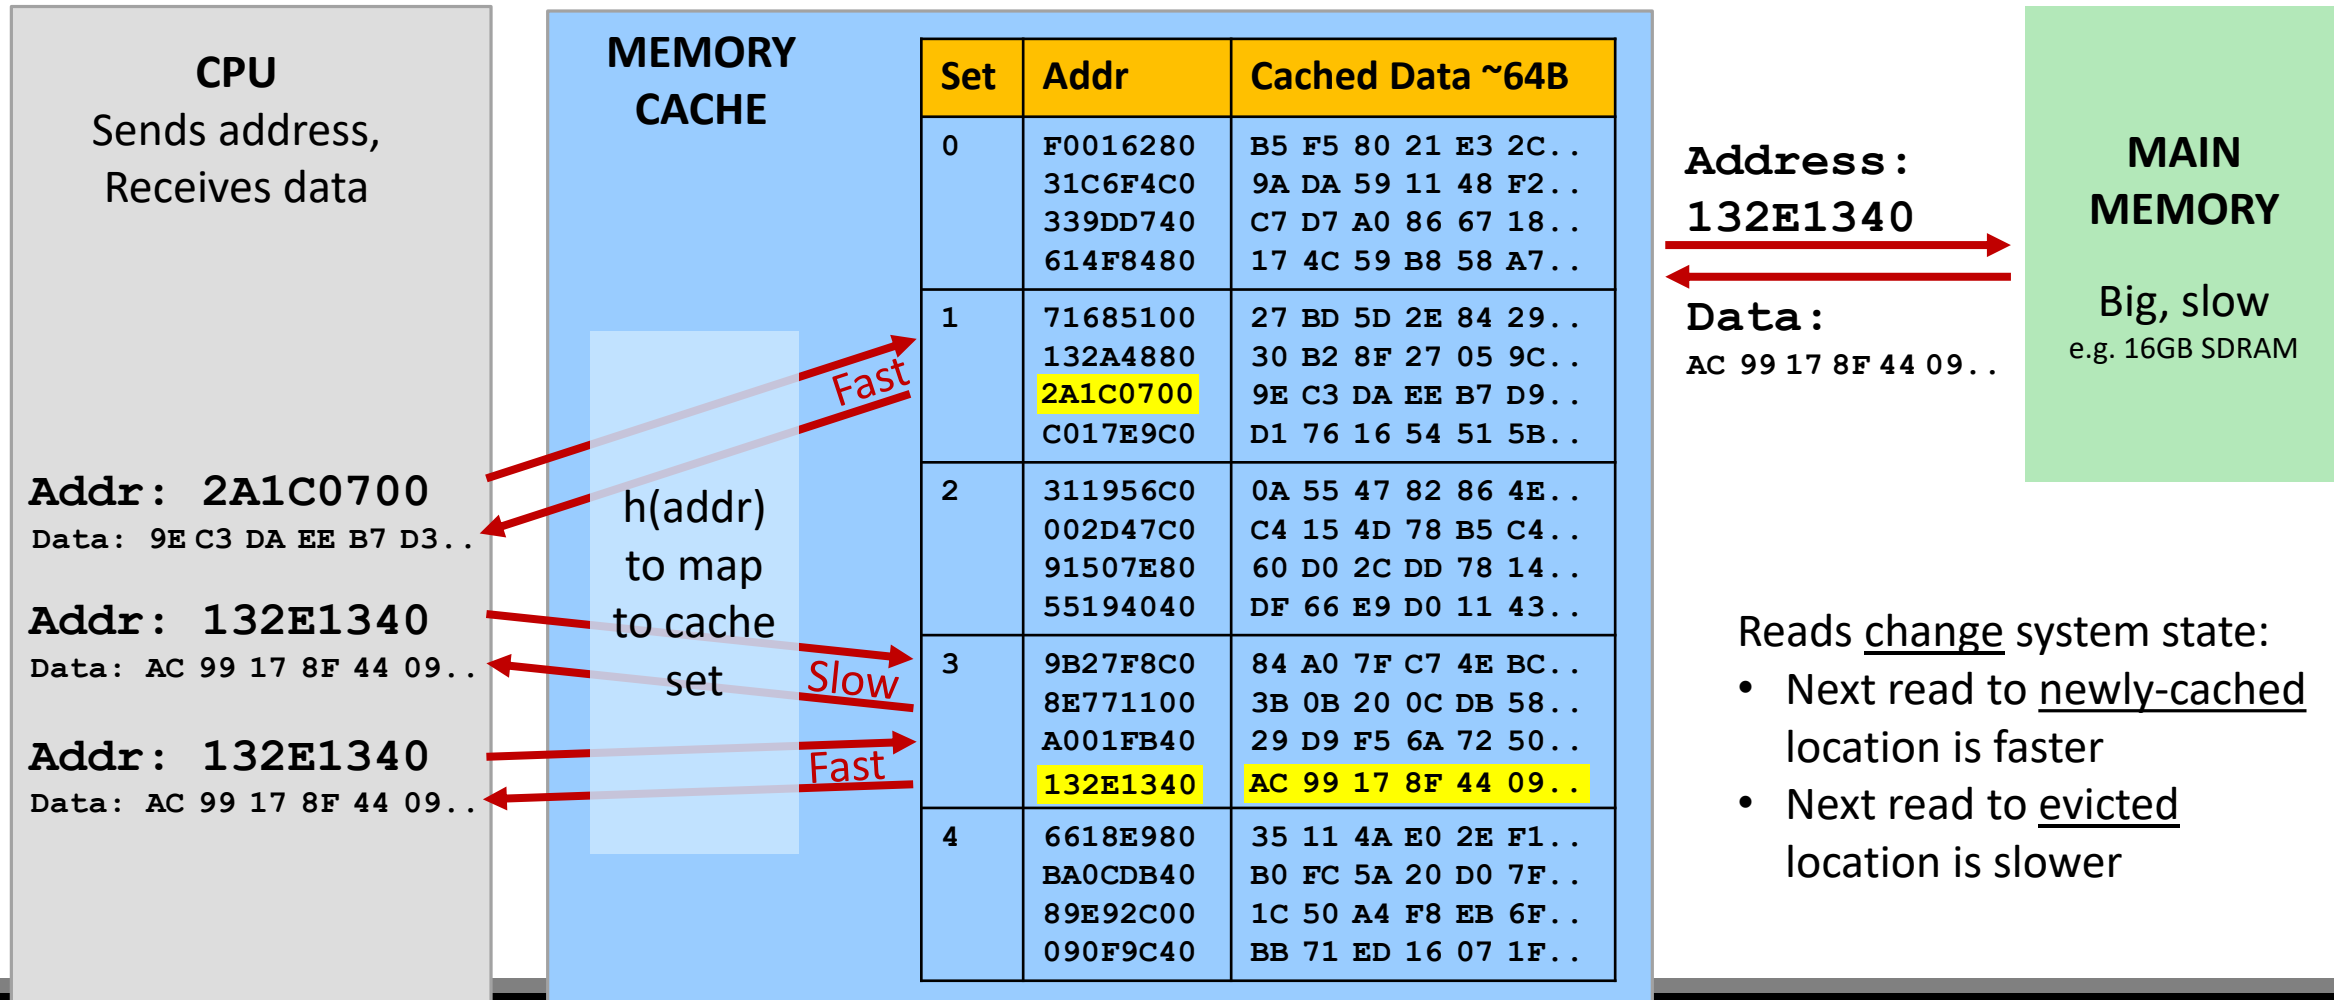
Optimization strategy

→ Caches = very fast accesses for frequently used memory locations

→ Work on multiple instructions at once (pipelining, speculation...)

Memory caches

Caches hold local (fast) copy of recently-accessed 64-byte chunks of memory



Cache attacks

Basic cache attack:

- Step 1: Prepare cache (erase or fill lines known data)
- Step 2: Run victim code
- Step 3: Detect changes in cache state (additions or removals)
- Step 4: Infer victim's secret data

Common variants include:

FLUSH+RELOAD

- Attacker flushes an address it shares with the victim (e.g. `clflush` instruction on an address in a shared library)
- Victim runs
- Attacker measure time to read from address (fast=victim accessed an address in the cache line)

PRIME+PROBE

- Attacker loads cache with attacker data, e.g. by reading data in attacker address space
- Victim runs
- Attacker reads back its data (slow=attacker data got evicted, e.g. because victim read data from an address mapping to same cache set)

EVICT+TIME

- See if evicting a location changes how long victim takes to run (e.g., run victim repeatedly, optionally evicting an address in between)

+ others (e.g., EVICT+RELOAD, FLUSH+FLUSH...)

Implications of shared caches:

- Adversary can infer the pattern of memory addresses accessed by adversary
- Adversary can infer the control flow, i.e. the destination of branch/jump instructions

Various responses/mitigations:

- Write code so secrets do not affect control flow & memory accesses [used for crypto libraries]
- Disable SMT (aka Hyperthreading)
- Page coloring, cache partitioning (e.g., Intel Cache Allocation Technology (CAT))
... but implementations have limits + some workarounds (e.g., see CATalyst) + other side channels
- Don't multitask attacker code on computers with sensitive data

Fault attacks

RSA Signing: $S = M^d \bmod n$

RSA is normally optimized using Chinese Remainder Theorem (faster)

- $S_p = (M \bmod p)^{d_p} \bmod p$
 - $S_q = (M \bmod q)^{d_q} \bmod q$
 - $S = S_q + q(k(S_p - S_q)M \bmod p)$
- M is padded message hash
d is secret exponent
n is public key (= p·q, where p and q are secret)

What if the CPU makes a mistake?

- Computation of S_p or S_q use most of the computation time, so are the most likely places for an error
- Example: Defective signature S' with erroneous S_p but correct S_q
 - Result: S' is wrong by a multiple of q
 - Attacker computes $\gcd(M - S'^e, n)$, which reveals q, which factors N

Fault attacks, continued

Not just RSA -- applicable to any crypto algorithm

- AES, 3DES, ECDSA...

Or bypass crypto altogether

- Example: Corrupt length entering output API

Various analog tricks to cause computation errors

- Pay TV glitchers: Send timed clock/ground glitches to smart card
- CLKSCREW: Misconfigure frequency/voltage management circuitry to induce errors
- Rowhammer: Corrupt data in DRAM via malicious memory access patterns
- Xbox 360 reset glitch attack
- Etc

Nobody would intentionally design a CPU that made errors... right?



Speculative execution

CPUs run instructions

- Correct result is defined as the result of performing instructions in-order

CPUs may run instructions out-of-order if this doesn't affect result

- Example:

```
a ← constant
b ← slow_to_obtain
c ← f(a)  // start before b finishes
```

CPUs can also *guess* likely program path and do speculative execution

- Example:

```
if (uncached_value_usually_1 == 1)
    compute_something()
```
- Branch predictor guesses that if() will be 'true' based on prior history
- Before uncached value known, save old registers/state and start executing compute_something() speculatively
- When memory read finishes, if() can be evaluated definitively:
 - Guess correct: Save speculative work – performance gain
 - Guess incorrect: Discard speculative work

The CPU is making erroneous computations –
effectively implementing a fault attack on the software



Conditional branch (Variant 1) attack

```
if (x < array1_size)
    y = array2[array1[x]*512];
```

Imagine this code was in a kernel API where `x` came from caller (e.g. attacker)

Execution without speculation is safe

- CPU will not evaluate `array2[array1[x]*512]` unless `x < array1_size`

What about with speculative execution?

Conditional branch (Variant 1) attack

```
if (x < array1_size)
    y = array2[array1[x]*512];
```

Attack setup:

- Train branch predictor to expect if() is true (e.g. call with `x < array1_size`)
- Evict `array1_size` and `array2[]` from cache

Memory & Cache Status

`array1_size = 00000008`

Memory at `array1` base address:

8 bytes of data (value doesn't matter)

[... lots of memory up to `array1` base+N...]

09 F1 98 CC 90 ... (something secret)

`array2[0*512]`
`array2[1*512]`
`array2[2*512]`
`array2[3*512]`
`array2[4*512]`
`array2[5*512]`
`array2[6*512]`
`array2[7*512]`
`array2[8*512]`
`array2[9*512]`
`array2[10*512]`
`array2[11*512]`
...

Contents don't matter
only care about cache **status**

Uncached

Cached

Conditional branch (Variant 1) attack

```
if (x < array1_size)
    y = array2[array1[x]*512];
```

Attacker calls victim code with $x=N$ (where $N > 8$)

- Speculative exec while waiting for `array1_size`
 - Predict that `if()` is true
 - Read address (`array1 base + x`) w/ out-of-bounds x
 - Read returns secret byte = **09** (fast – in cache)

Memory & Cache Status

`array1_size = 00000008`

Memory at `array1` base address:

8 bytes of data (value doesn't matter)

[... lots of memory up to `array1 base+N...`]

09 F1 98 CC 90 ... (something secret)

`array2[0*512]`
`array2[1*512]`
`array2[2*512]`
`array2[3*512]`
`array2[4*512]`
`array2[5*512]`
`array2[6*512]`
`array2[7*512]`
`array2[8*512]`
`array2[9*512]`
`array2[10*512]`
`array2[11*512]`
...

Contents don't matter
only care about cache **status**

Uncached

Cached

Conditional branch (Variant 1) attack

```
if (x < array1_size)
    y = array2[array1[x]*512];
```

Attacker calls victim code with $x=N$ (where $N > 8$)

- Speculative exec while waiting for `array1_size`
 - Predict that `if()` is true
 - Read address (`array1 base + x`) w/ out-of-bounds x
 - Read returns secret byte = **09** (fast – in cache)
 - Request memory at (`array2 base + 09*512`)
 - Brings `array2[09*512]` into the cache
 - Realize `if()` is false: discard speculative work
- Finish operation & return to caller

Attacker times reads from `array2[i*512]`

- Read for $i=09$ is fast (cached), revealing secret byte

Memory & Cache Status

`array1_size = 00000008`

Memory at `array1` base address:

8 bytes of data (value doesn't matter)

[... lots of memory up to `array1 base+N...`]

09 F1 98 CC 90 ... (something secret)

`array2[0*512]`
`array2[1*512]`
`array2[2*512]`
`array2[3*512]`
`array2[4*512]`
`array2[5*512]`
`array2[6*512]`
`array2[7*512]`
`array2[8*512]`
`array2[9*512]`
`array2[10*512]`
`array2[11*512]`
...

Contents don't matter
only care about cache **status**

Uncached

Cached

Violating JavaScript's sandbox

JavaScript code runs in a sandbox

- Not permitted to read arbitrary memory
- No pointers, array accesses are bounds checked

Browser runs JavaScript from untrusted websites

- JavaScript engine can interpret code (slow) or compile it (JIT) to run faster
- In all cases, engine must enforce sandbox (e.g. apply bounds checks)

Speculative execution can blast through safety checks...

- Can we write JavaScript that compiles into machine code that leaks memory contents?

Violating JavaScript's sandbox

index will be in-bounds on training passes, and out-of-bounds on attack passes

JIT thinks this check ensures **index** < **length**, so it omits bounds check in next line. Separate code evicts **length** for attack passes

```
if (index < simpleByteArray.length) {  
  index = simpleByteArray[index | 0];  
  index = (((index * TABLE1_STRIDE) | 0) & (TABLE1_BYTES - 1)) | 0;  
  localJunk ^= probeTable[index | 0] | 0;  
}
```

Do the out-of-bounds read on attack passes!

"|0" is a JS optimizer trick (makes result an integer)

4096 bytes (= page size)

Need to use the result so the operations aren't optimized away

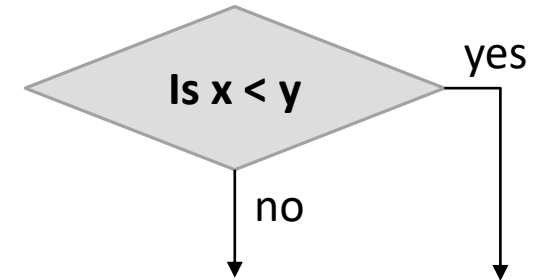
Keeps the JIT from adding unwanted bounds checks on the next line

Leak out-of-bounds read result into cache state!

Indirect branches

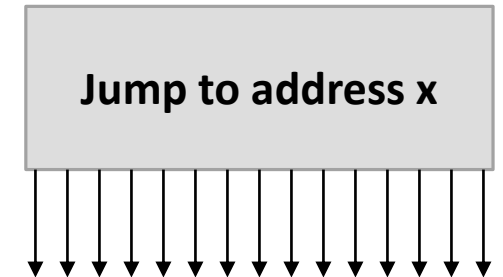
Conditional branches: Fork in the road

- E.g. next instruction if false, jump destination if true



Indirect branches: Teleport anywhere

- Examples on x86: `jmp [1234567]` `jmp eax` `ret`
- Branch target buffer: CPU tracks past jump destinations to make quick guesses
- If destination is delayed, CPU guesses and proceeds speculatively



Vastly more freedom for attacker – billions of possible destinations

Poisoning indirect branches

Pick indirect branch to redirect speculatively

- E.g. a jump that occurs with attacker-controlled values in some registers

Pick destination for victim to speculatively execute (the “gadget”)

- Instructions in victim code/libraries usable to leak victim’s memory into covert channel
- Like return oriented programming, but gadget also doesn’t have to return nicely

Attack

- **Mistrain** branch prediction/BTB so speculative execution will go to gadget
- **Evict** or flush destination address from the cache to ensure long duration speculative execution
- **Execute** victim so it runs gadget speculatively
- **Detect** change in cache state (e.g. EVICT+RELOAD) to determine memory data
- **Repeat** for more bytes

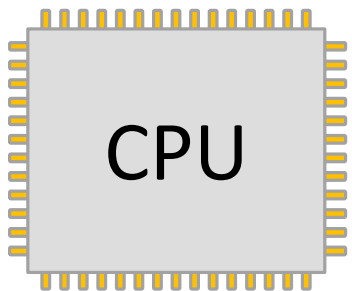
Mitigations

Mitigation. noun. “The action of reducing the severity, seriousness, or painfulness of something”

Not necessarily a complete solution

Public domain image by Svetlana Miljkovic
Definition from <https://en.oxforddictionaries.com/definition/mitigation>





Variant 1 Mitigation: Speculation-barrier instruction (e.g. LFENCE)

- Idea: Software developers insert barrier on all vulnerable code paths in software
- Simple & effective (for CPU developer)



CPU architecture
Machine language
Compiler
Higher-level language

Abstraction boundaries

Software

- operating systems
- drivers
- web servers
- interpreters/JITs
- databases
- :



Insert LFENCES manually?

Often millions of control flow paths

Too confusing - speculation runs 188++ instructions, crosses modules

Too risky – miss one and attacker can read entire process memory

Put LFENCES everywhere?

Abysmal performance - LFENCE is very slow (+ no tools)

Not in binary libraries, compiler-created code patterns

Insert by smart compiler?

Protect all potentially-exploitable patterns = too slow

- Compilers judged by performance, not security

Protect only known-bad bad patterns = unsafe

- Microsoft Visual C/C++ /Qspectre unsafe for 13 of 15 tests

<https://www.paulkocher.com/doc/MicrosoftCompilerSpectreMitigation.html>

Unstable: Effects are intentionally hidden by CPU + can differ in new CPUs

- Transfers blame to software developer (“you should have put an LFENCE there”)

Mitigations: Indirect branch variant

x86:

- New MSRs created via microcode
 - Low-level control over branch target buffer
 - Intel + AMD: Microcode updates for many (but not all) CPUs: messy (stability problems + recalls...)
 - Available from kernel mode only (use by applications unclear)
 - Performance impact
- Retpoline proposal from Google
 - Messy hack -- replace indirect jumps with construction that resists indirect branch poisoning on Haswell
 - Microcode updates to make retpoline safe on Skylake & beyond
- **Really hard to test**

“All of this is pure garbage”
-- Linux Torvalds
<https://lkml.org/lkml/2018/1/21/192>

ARM: No generic mitigation option

- Fast ARM CPUs broadly impacted, e.g. Cortex-A9, A15, A17, A57, A72, A73, A75...
- Often no mitigation, but on some chips software may be able to invalidate/disable branch predictor (with “non-trivial performance impact”)
 - See: <https://developer.arm.com/support/security-update/download-the-whitepaper>

DOOM with only MOV instructions

Only MOV instructions

- No branch instructions
- One big loop with exception at the end to trigger restart

Sub-optimal performance

- One frame every ~7 hours

A branchless DOOM

This directory provides a branchless, mov-only version of the classic DOOM video game.



DOOM, running with only mov instructions.

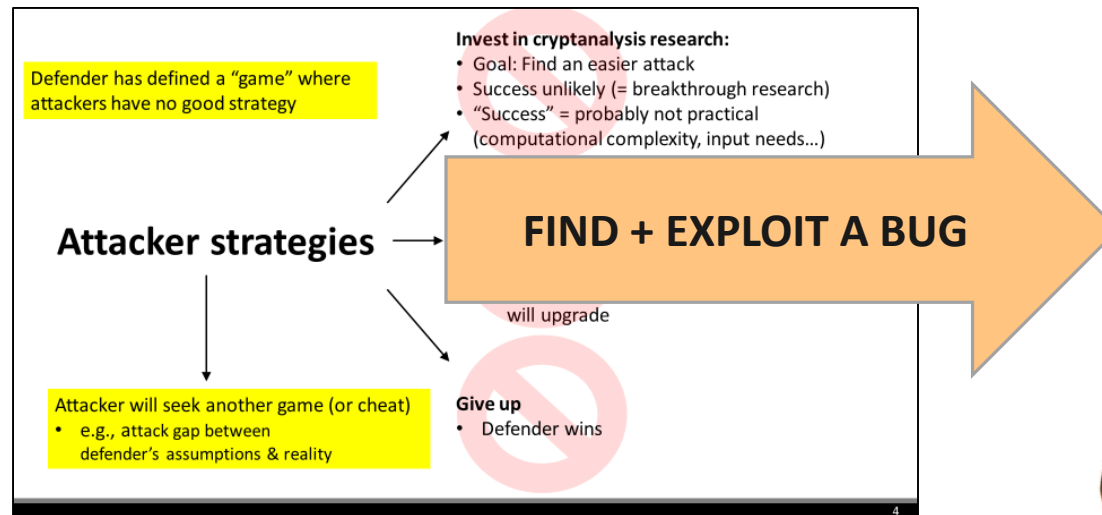
This is thought to be entirely secure against the Meltdown and Spectre CPU vulnerabilities, which require speculative execution on branch instructions.

<https://github.com/xoreaxeaxeax/movfuscator/tree/master/validation/doom>

Implications + Looking to the Future

Risk in context

Because of software bugs, computer security was in a dire situation



Spectre doesn't change the magnitude of the risk, but adds to the mess
Complexity of fixes -> new risks
Psychology of unfixed vulnerabilities

Kerckhoffs's principle (1883) :

A system should be secure even if everything about it, except the key, is public knowledge.

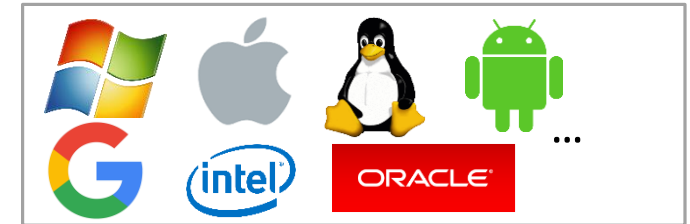
1. Insecure: The designer knows the security is weak

$$P(\text{secure}) = 0$$



2. Ignorance: Attackers can almost certainly succeed using what designers don't know

$$P(\text{secure}) < \epsilon \quad (\text{updates can maintain security at designer ignorance})$$



3. Probably secure: Reasonable chance of no bugs

$$50\% < P(\text{secure}) < 99+\%$$



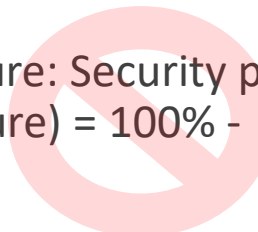
AES, seL4, TLS 1.2+(?), simple hardware designs??

If the assumptions correct... ☹️

4. Reliably secure: Security properties which anyone can easily verify and are widely checked

$$P(\text{secure}) = 100\% - \epsilon$$

Hopelessly difficult?



Reaction



At Least Three Billion Computer Chips Have the Spectre Security Hole

Companies are rushing out software fixes for Chipmageddon.

Intel Faces Scrutiny as Questions Swirl Over Chip Security

Silicon melts

Spectre and Meltdown prompt tech industry soul-searching

"AMD is not susceptible to all three variants. [...] there is a near zero risk to AMD processors at this time."

Researchers Discover Two Major Flaws in the World's Computers

ANDY GREENBERG SECURITY 01.03.18 03:00 PM

A CRITICAL INTEL FLAW BREAKS BASIC SECURITY FOR MOST COMPUTERS

Apple says Spectre and Meltdown vulnerabilities affect all Mac and iOS devices

??!

Why panic when a vulnerability is identified?

$P(\text{secure})$ fell from ϵ to 0.

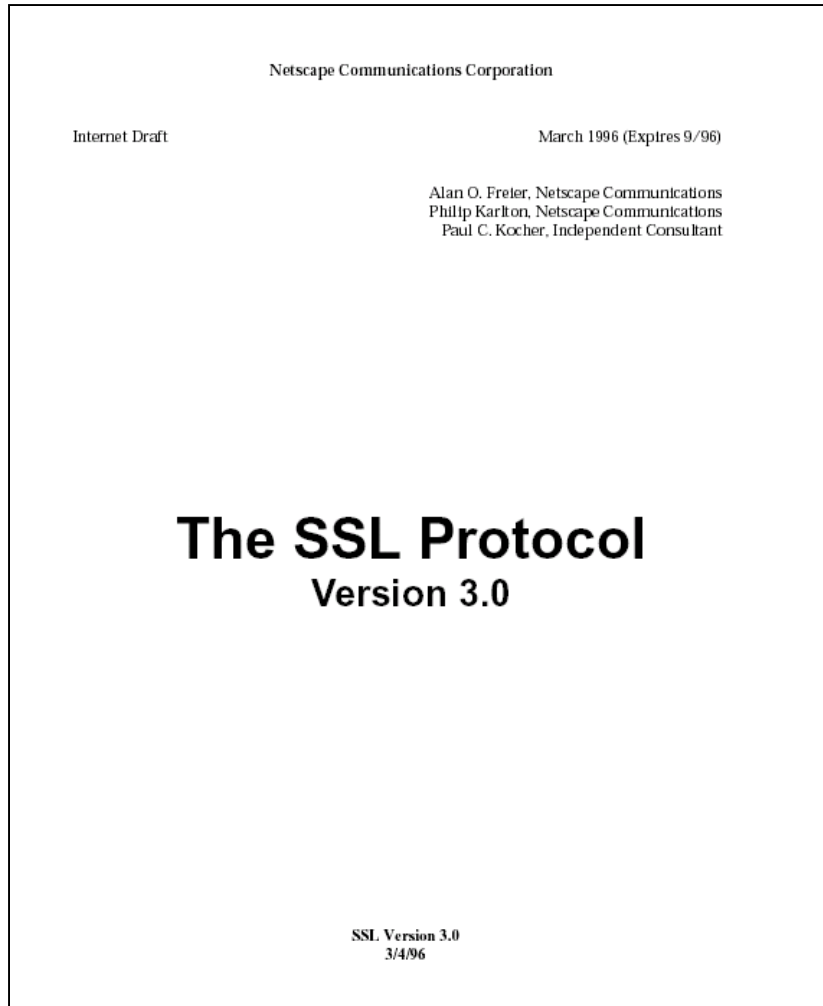
but ϵ was usually negligible

Optimist's security = $\lceil P(\text{secure}) \rceil$

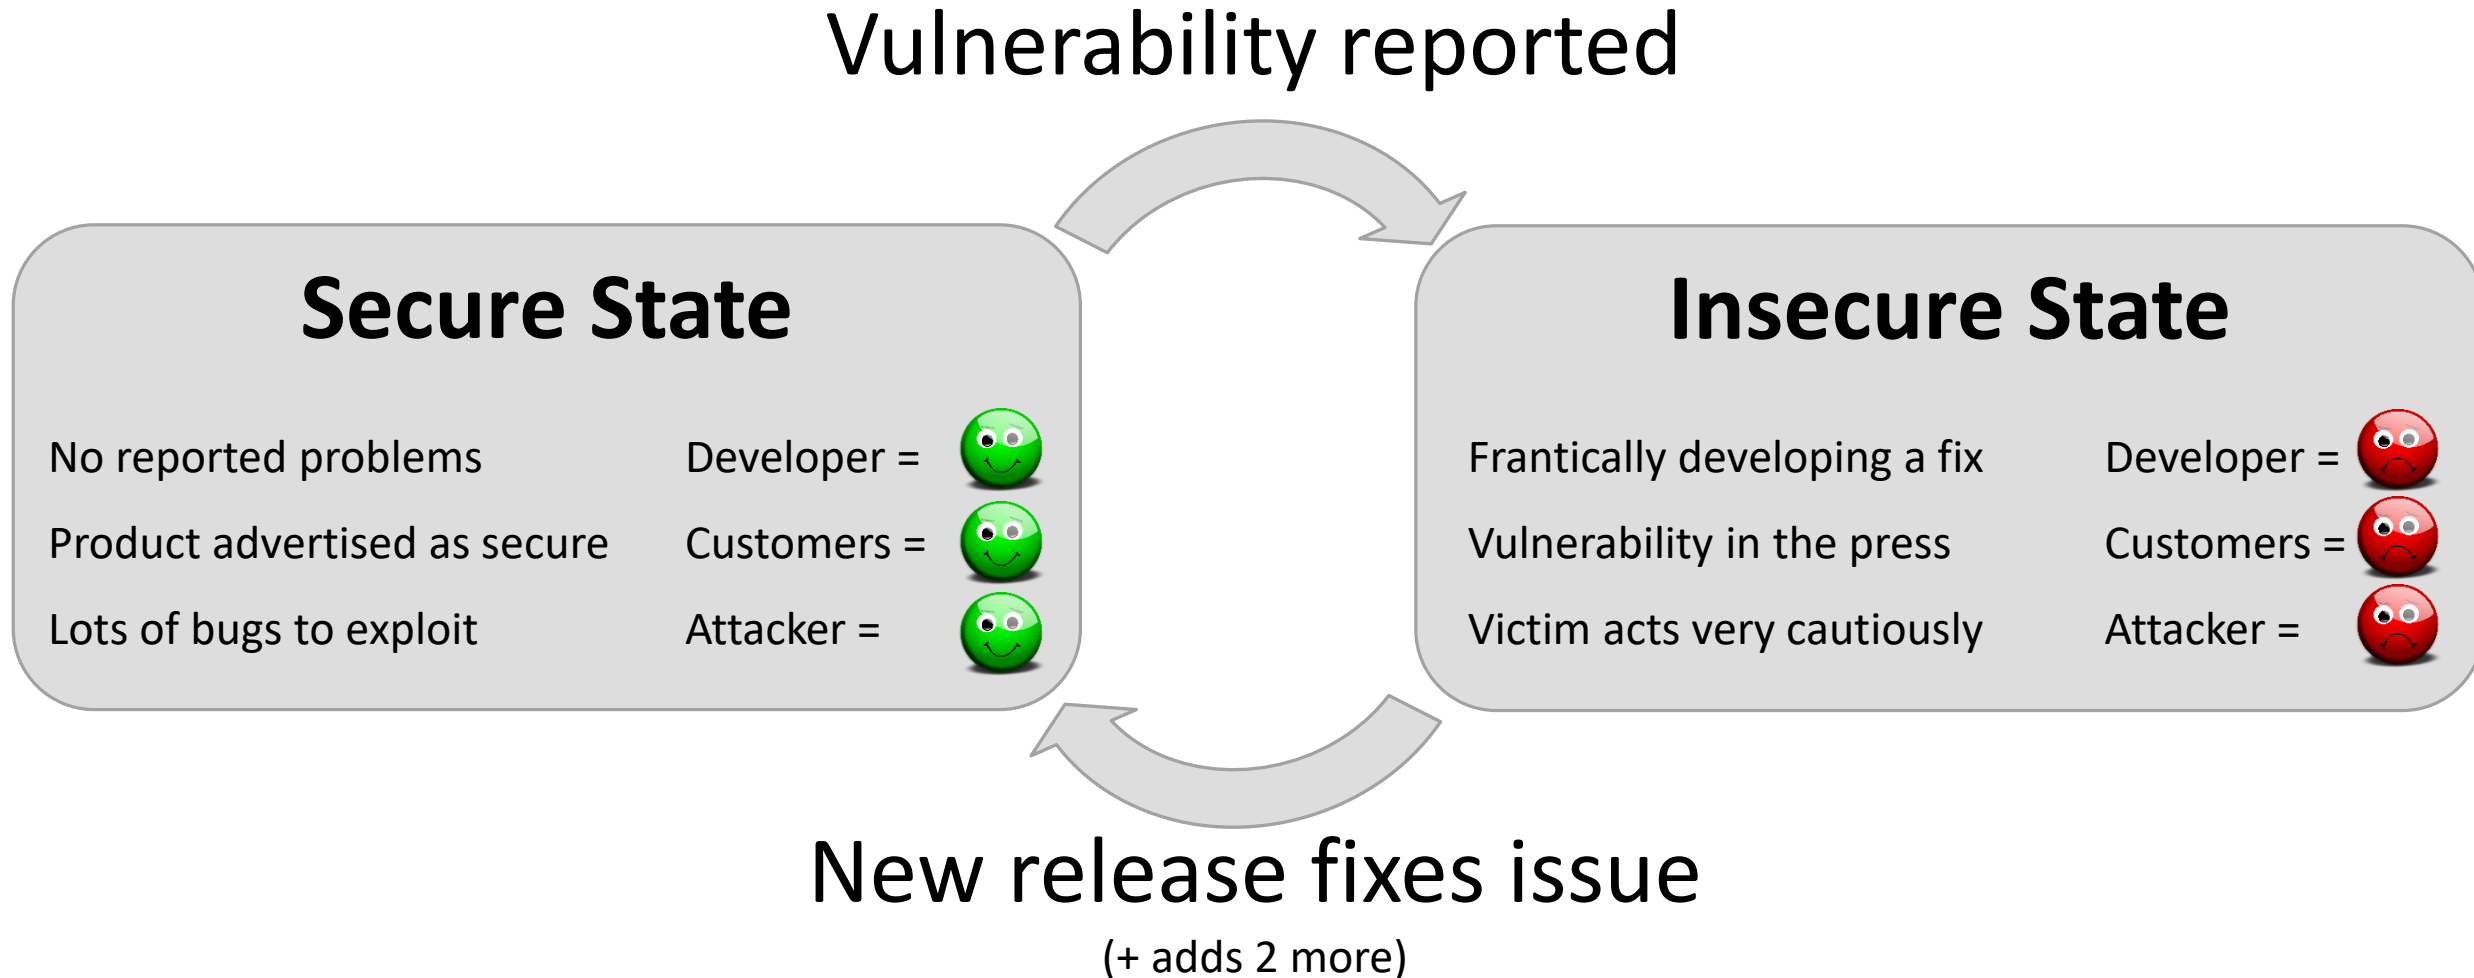
fell from 100% to 0%



Long history of over-optimism



State machine for the security cycle



A great question

Are there any security implications
from speculative execution?

-- Mike Hamburg

at the intersection of two trends



Semiconductor/CPU evolution

The diagram consists of two large, light-gray arrows pointing to the right. The top arrow is wider and contains the text 'Semiconductor/CPU evolution'. The bottom arrow is narrower and contains the text 'Security'. The two arrows overlap, with the 'Security' arrow positioned slightly below and to the right of the 'Semiconductor/CPU evolution' arrow, creating a visual representation of their intersection.

Security

Why a great question?

Are there any security implications from speculative execution?

-- *Mike Hamburg*

Explores gaps...

Architecture vs. CPU

- Programmers' mental model: simple in-order CPUs
- Reality: Enormous hidden complexity of actual CPUs

Correctness assumptions

- Security efforts rely on CPUs executing code faithfully
- Reality: Actual CPUs are making then discarding errors

Expertise

- Few security experts specialize in CPU details
- Few CPU designers specialize in security

Hidden in plain sight

Obvious... in hindsight:

- Component behaviors well known
(taught in textbooks on electronics/CPU design)

Why missed until now?

- Specialization of technical teams
 - Implications (or existence) of properties not obvious to others
- Few communications channels, different terminology, perspectives, assumptions, goals
 - Hardware designers <-> software developers
 - Software developers <-> security specialists
 - Security specialists <-> cryptographers

} Focus on minutiae (bugs, individual blocks) rather than risks
Few cross-disciplinary people, conferences...
- Assumption that old approaches meet current changing requirements
- Failure to explore implications of prior side channels results
- Performance-first priorities

“The reality is there are probably other things out there like [Spectre] that have been deemed safe for years. Somebody whose mind is sufficiently warped toward thinking about security threats may find other ways to exploit systems which had otherwise been considered completely safe.”

– Simon Segars (CEO, Arm Holdings)

Are there any security implications
from speculative execution?

-- Mike Hamburg

```
int max=100;  
int array1[100]  
int array2[.];
```

Regent
TAIPEI

Tuval
Daniel Geler
mike H

x = value from attacker

flush (add)

```
if (x < 0 || x > max)  
    return (-1);
```

```
byte y = array1[x]  
z = array2[y].fst  
return (z);
```

- ① clear cache
- ② bang key into cache
- ③ call this.

+ asked before an uninteresting
conference talk

What are some other areas we could ask questions about?

Confidentiality (leaking data)

- Side channel attacks: Spectre variants + AI/ML analysis...
- Analog effects (voltage sensors, PLLs, RF leaks to radios...)
- Incomplete zeroization on context switch (debug regs...)

Integrity (corrupting information)

- Crosstalk effects in chips
- Clock control logic
- Misconfigure power mgmt/SW-induced glitches
- Rare failures on big distributed systems

Availability (crashing/destroying hardware)

- NVMs fail after 10K-1M writes - is write-leveling secure?
- Overheating (disable thermal shutoff + cause heating)
- Electromigration & other physical effects
- Malicious blowing of (anti-)fuses
- FPGA failures after to loading malicious bitstreams

Manufacturing, infrastructure, legal

- Insecure test modes (global JTAG passwords common)
- Insecure factory key programming
- Political/legal threats (complex mandatory backdoors...)
- Attacks on firmware/microcode signing keys

Can we better quantify risk?

Many aspects: data quality, analysis cost, reproducibility, incentives:

- Vendors – want to support/fund work that highlights security (not risk)?
- Customers – lack access to needed data, uncoordinated spending, limited choices... and are optimists!
- Regulators?

Security Facts

DANGER: This product is not secure
Probability of critical defects >99%

Complexity measures

| | |
|-------------------------|-------------------------|
| Total complexity | 2190K LoC |
| Security-critical areas | 791K LoC |
| Bug density in critical | 3×10^{-4} /LoC |

Risk estimates

| | |
|---------------------|-------------|
| Remote exploit flaw | 80% |
| Local non-invasive | >99% |
| Local invasive | unprotected |

Duration of security support: 10 yr

After 10 years use prohibited

Security Facts

This product is probably secure

Complexity measures

| | |
|---|-------------------------|
| Total complexity | 1781K LoC |
| Security-critical areas | 37K LoC (x2*) |
| * Full hardware and software redundancy | |
| Bug density in critical | 2×10^{-4} /LoC |

Risk estimates

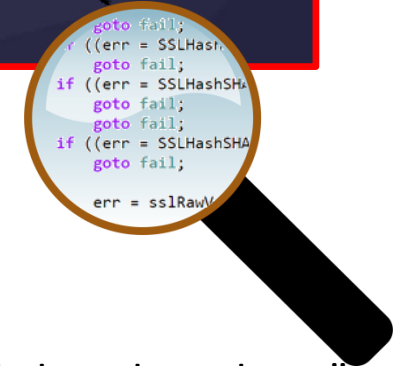
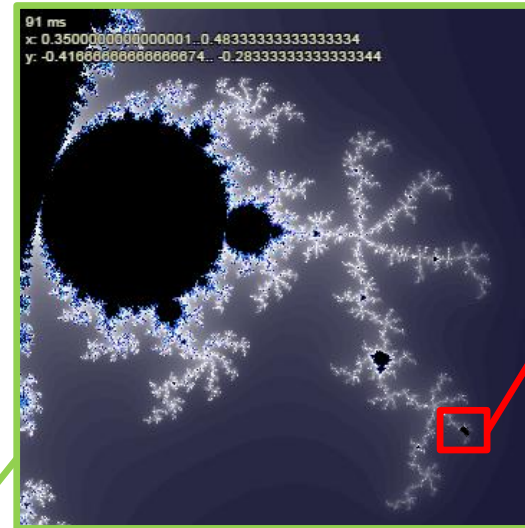
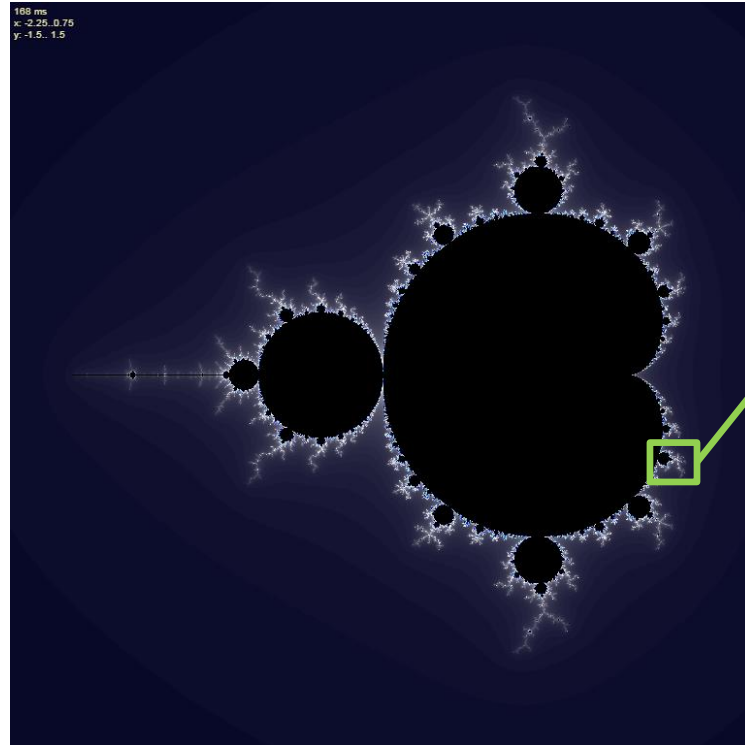
| | |
|---------------------|-------------|
| Remote exploit flaw | <3% |
| Local non-invasive | <15% |
| Local invasive | unprotected |

Duration of security support: 25 yr

After 25 years open source

New & Improved!
Now only a 10% chance of
being completely hackable.

Problems = Opportunities



Overall risks are “obvious”
– “Software and hardware are very likely to have bugs”

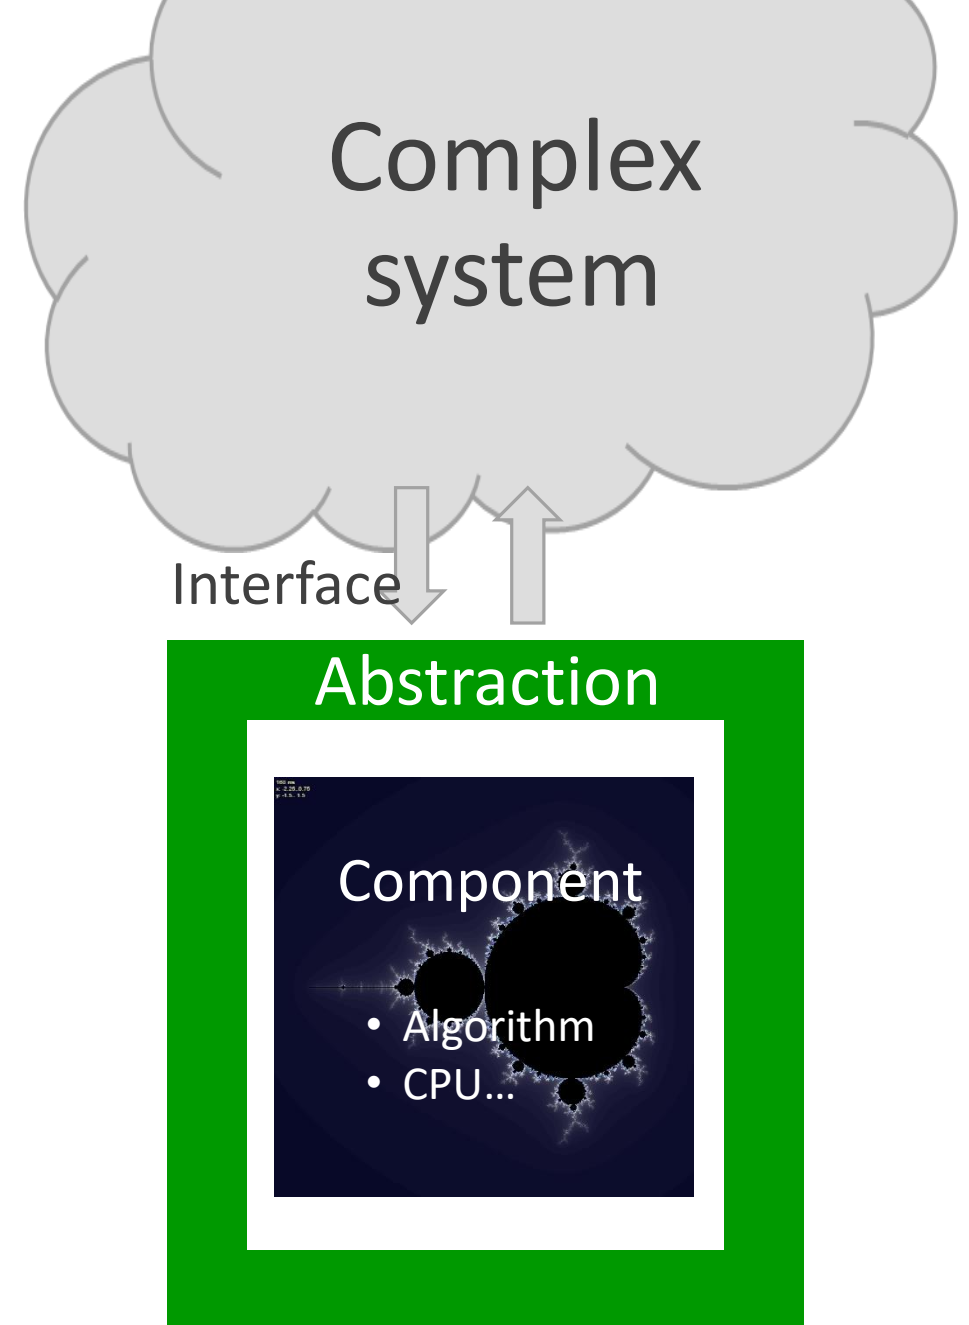
Individual bugs are “obvious”
– when we stare directly at them

What can we do to reduce the obvious risks without
having to understand every logic gate & line of code?

The vulnerabilities in this talk all arise from critical details of a component not described by abstraction compromising the security a complex system.

We need 4 things:

- Better components
 - = Better matched to system security needs
- Better abstraction/interface descriptions
 - = Conveys security properties
- Better ways to instantiate components
 - = Use components with less risk to overall system
- Better reasoning about compositions
 - = Understand combinations of components



Strategies from yesterday's talk

- Abstraction collapsing
- Temporal separation
- Spatial separation
- Redundancy

Software determinism

- Can we guarantee (most) SW is deterministic?
 - Output = $f(\text{source}, \text{input})$... and nothing else
 - Functional programming attributes for large functional blocks
 - Repeatable, checkable
- Can use for a lot (but not everything)
 - Codecs, decompression, font rendering, math...
 - Much lower overhead (+ better compartmentalization) than microservices
- Strategies for realizing
 - **sound** static analyzers
 - safe subsets of unsafe languages (C, assembly)
 - safer compilers
 - safer languages

Implies:

- Same output on all CPUs
- No buffer overflows
- No race conditions
- No side channel receivers
... but can transmit

Note: Determinism \neq correctness
Narrower but more achievable

Challenge: Poor hardware support!

Example: Big performance cost for dynamic checks due to lack of support

We need better hardware!

In essentially all practical situations, computations can only be as secure as the hardware that performs them

- Technical: Better hardware can address technical barriers to safer systems & safer software. [Example: CHERI]
- Stronger foundational layers (e.g. HW) is needed to make security investments in higher layers make economic sense

No “best” architecture or implementation

Fastest != safest

- Market (mis-)trained to expect optimal everything
- Result: same CPUs for video games, funds transfers, ...

Trade-offs: Need to bifurcate faster vs. safer

- ‘Safer’ needs to be much less complex HW and OS
-- not just a different mode (like TrustZone/SGX)
- Can co-exist in one device (or chip)



Many areas of difference:

- performance needs
- tolerance for complexity-induced risks
- choice of risky optimizations
- safety margins
- side channel/fault attack countermeasures
- backward compatibility with unsafe software
- feature choices
- verification/test processes
- ...



Volvo image public domain by OSX, source: [https://en.wikipedia.org/wiki/Volvo_XC90#/media/File:2005_Volvo_XC90_\(P28_MY05\)_2.5_T_wagon_\(2011-11-18\)_01.jpg](https://en.wikipedia.org/wiki/Volvo_XC90#/media/File:2005_Volvo_XC90_(P28_MY05)_2.5_T_wagon_(2011-11-18)_01.jpg)
Race car by Andrew Griffith CC BY 2.0, source: https://en.wikipedia.org/wiki/Ferrari_F10#/media/File:Felipe_Massa_Ferrari_Bahrain_2010_GP.jpg
SR-71 image in public domain, source: [https://commons.wikimedia.org/wiki/File:Lockheed_SR-71_Blackbird_\(modified\).jpg](https://commons.wikimedia.org/wiki/File:Lockheed_SR-71_Blackbird_(modified).jpg)
787 image by Dave Sizer, CC BY 2.0, source: https://simple.wikipedia.org/wiki/Boeing_787_Dreamliner#/media/File:787_First_Flight.jpg

10 areas for work

① Side channel attacks

- What analog & digital effects leak information?
- What new optimizations will leak information?

② Fault attacks

- How might errors be induced (crosstalk, glitches...)? Detected? Mitigated?

③ Denial of service attacks

- Can software induce permanent failures (e.g. damage hardware via NVM exhaustion, fuses, electromigration...)?

④ Architectures

- What guarantees “should” architectures make to meet the needs of SW? How should guarantees be documented? How do guarantees flow to code in higher level languages?

⑤ Memory models

- How should memory work in secure systems? What support is needed in higher-level languages & legacy code?

⑥ Verification & test

- How to minimize (or just estimate) the likelihood of vulnerabilities in architectures? Implementations? Manufactured devices?
- Are there safer ways to test manufactured devices (e.g. scan chains can enable attacks)?

⑦ Metrics for security

- What metrics might help buyers, developers, and regulators make informed security decisions?

⑧ Legacy migration

- How can existing architectures (x86, ARM), operating systems, drivers, and applications be migrated to something safer?

⑨ Incident response

- How should hardware vulnerabilities be handled, including in cases where devices cannot be fixed?

⑩ Education

- How can we meet the need for security talent (research, architecture, implementation, verification, operation...)?

Q&A

My email: paul@paulkocher.com